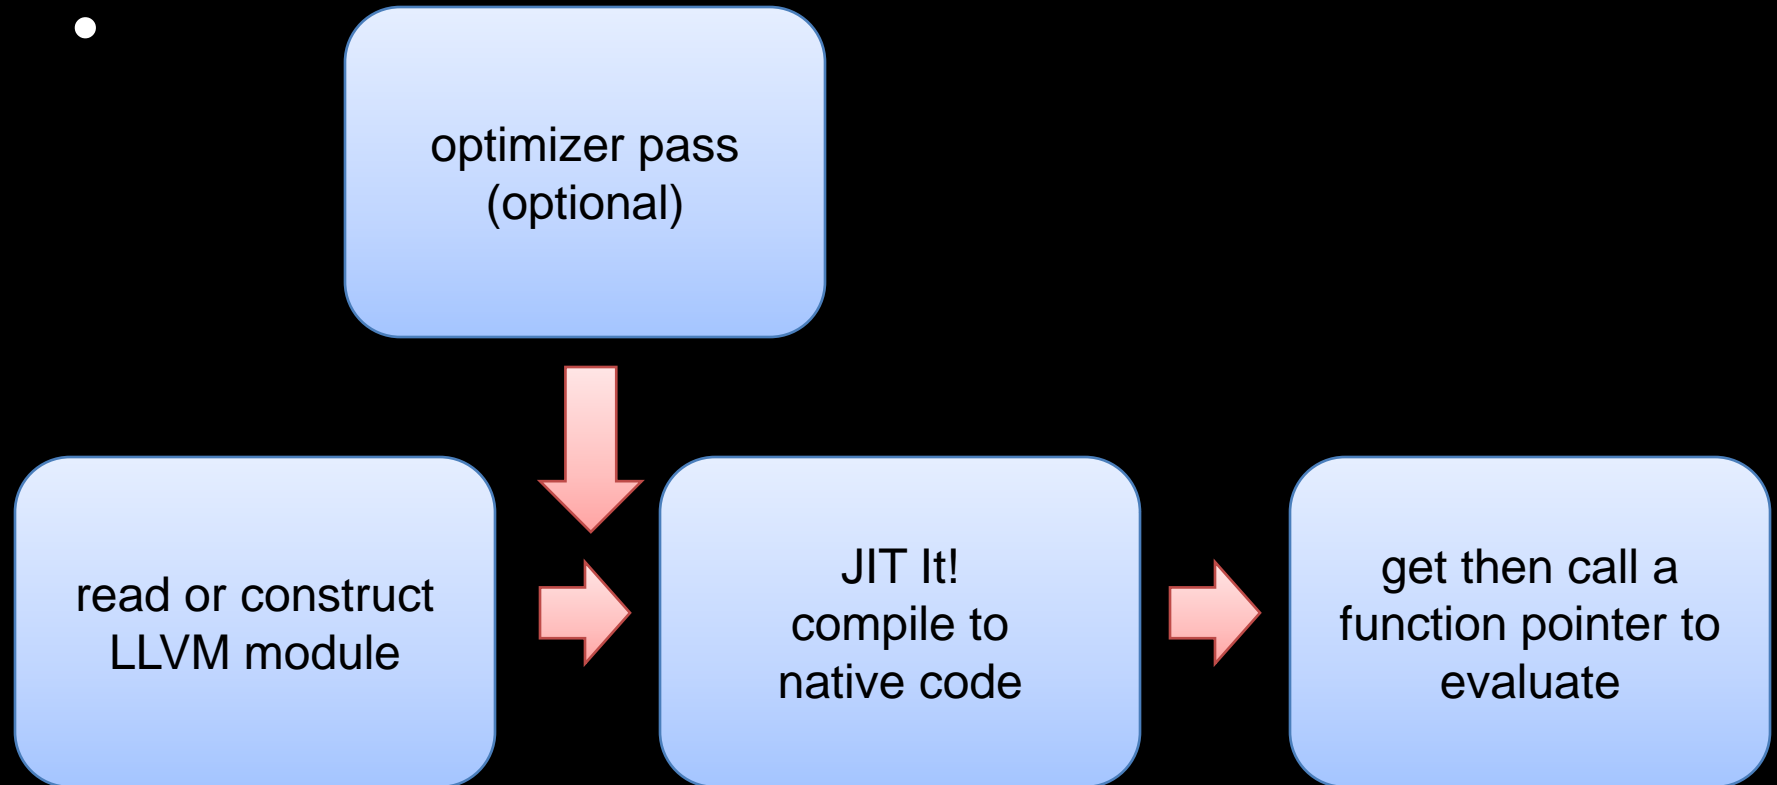


How To Use JIT & Make JIT Support

Agenda

- How to use JIT
 - Construct Module
 - Read Module
 - JIT & Execute
- Make JIT support

How to use JIT



What is LLVM module

- LLVM programs are composed of "Module"s
- a translation unit of the input programs
- consists of functions, global variables, and symbol table entries

"hello world" module

; Declare the string constant as a global constant...

```
@.LC0 = internal constant [13 x i8] c"hello world\0A\00" ;
```

; External declaration of the puts function

```
declare i32 @puts(i8 *) ; i32(i8 *)*
```

; Definition of main function

```
define i32 @main() {
```

```
    ; Convert [13x i8]* to i8*...
```

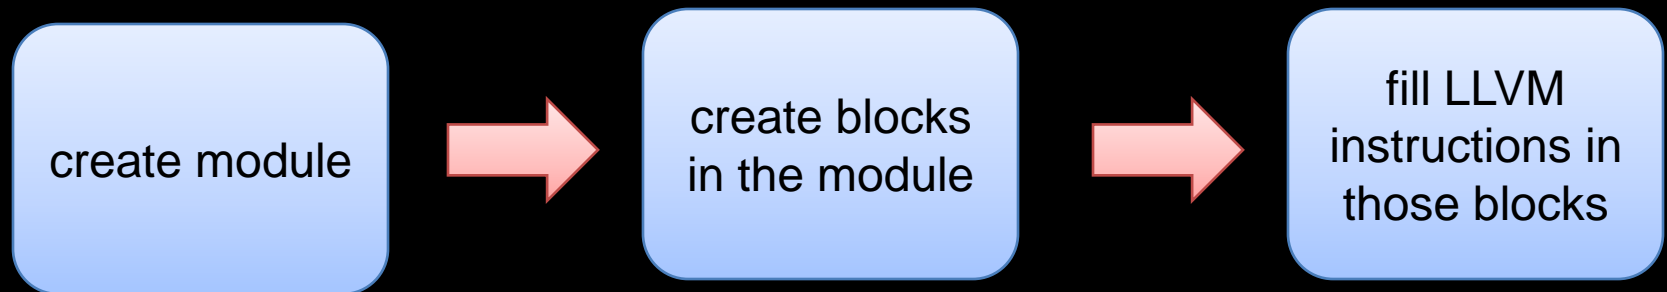
```
    %cast210 = getelementptr [13 x i8]* @.LC0, i64 0, i64 0
```

```
    ; Call puts function to write out the string to stdout...
```

```
    call i32 @puts(i8* %cast210) ; i32 ret i32 0
```

```
}
```


Construct LLVM module



Greatest Common Divisor (GCD) algorithm

```
unsigned gcd( unsigned x, unsigned y ) {  
  
    if ( x == y ) {  
        return x;  
    } else if ( x < y ) {  
        return gcd( x, y - x );  
    } else {  
        return gcd( x - y, y );  
    }  
  
}
```

CFG for 'gcd' function



LLVM IR

```
; ModuleID = 'tut2'

define i32 @gcd(i32 %x, i32 %y) {
entry:
    %tmp = icmp eq i32 %x, %y                ; <i1> [#uses=1]
    br i1 %tmp, label %return, label %cond_false

return:
    ; preds = %entry
    ret i32 %x

cond_false:
    ; preds = %entry
    %tmp2 = icmp ult i32 %x, %y              ; <i1> [#uses=1]
    br i1 %tmp2, label %cond_true, label %cond_false1

cond_true:
    ; preds = %cond_false
    %tmp3 = sub i32 %y, %x                   ; <i32> [#uses=1]
    %tmp4 = call i32 @gcd(i32 %x, i32 %tmp3) ; <i32> [#uses=1]
    ret i32 %tmp4

cond_false1:
    ; preds = %cond_false
    %tmp5 = sub i32 %x, %y                   ; <i32> [#uses=1]
    %tmp6 = call i32 @gcd(i32 %tmp5, i32 %y) ; <i32> [#uses=1]
    ret i32 %tmp6
}
```

#include the appropriate LLVM header files

```
#include <llvm/Module.h>
```

```
#include <llvm/Function.h>
```

```
#include <llvm/PassManager.h>
```

```
#include <llvm/Analysis/Verifier.h>
```

```
#include <llvm/Assembly/PrintModulePass.h>
```

```
#include <llvm/Support/IRBuilder.h>
```

```
using namespace llvm;
```

main()

```
Module* makeLLVMModule(); // do the real work of creating the module
```

```
int main(int argc, char**argv) {
```

```
    Module* Mod = makeLLVMModule(); // creating a module
```

```
    verifyModule(*Mod, PrintMessageAction); // verifying it
```

```
    PassManager PM;
```

```
    PM.add(new PrintModulePass(&llvm::cout));
```

```
    PM.run(*Mod); // running the PrintModulePass on it
```

```
    delete Mod;
```

```
    return 0;
```

```
}
```

verifyModule()


- print an error message if your LLVM module is malformed

```
%x = add i32 1, %x
```

- the definition of %x does not dominate all of its uses, **NOT** a legal SSA form (every variable is assigned exactly once)

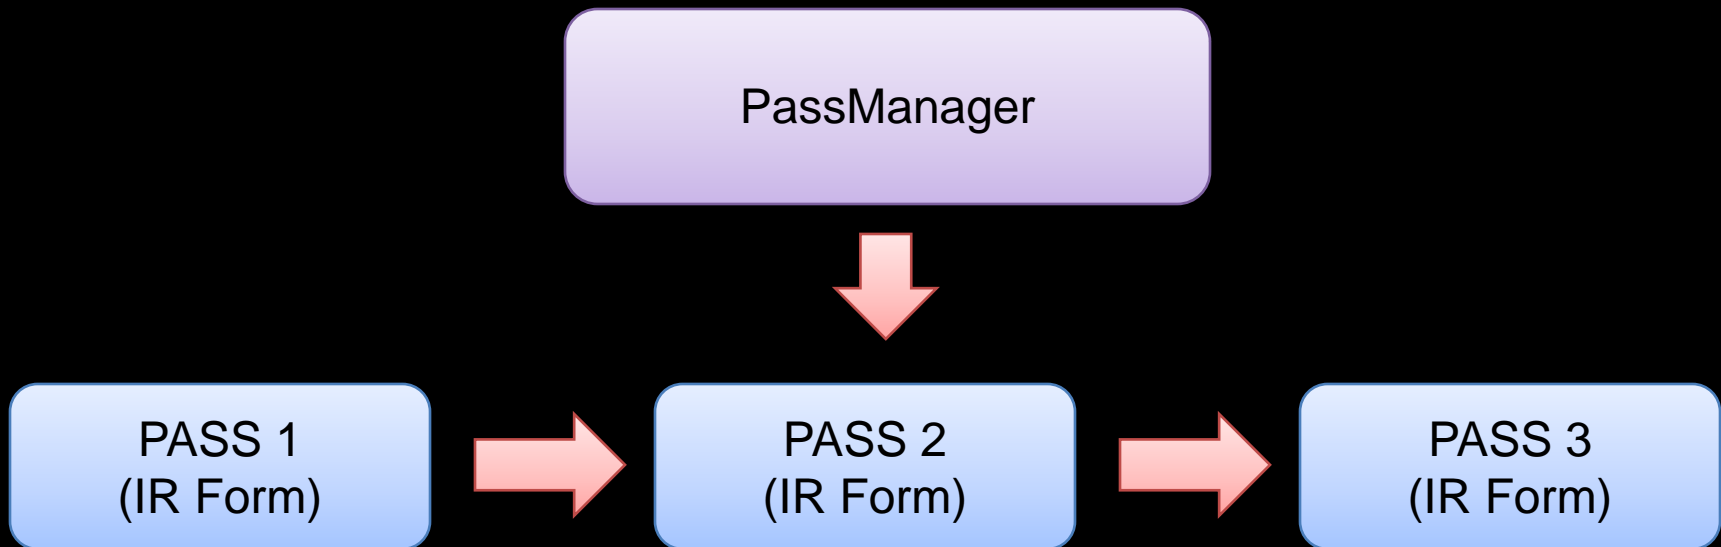
What is a pass?

- perform the transformations and optimizations that make up the compiler
- build the analysis results that are used by these transformations



What PassManager does

- takes a list of passes, ensures their prerequisites are set up correctly, and then schedules passes to run efficiently



PrintModulePass

- a pass that prints out the entire module when it is executed

```
PM.add(new PrintModulePass (&llvm::cout) );
```

- add this pass and schedule it

makeLLVMModule()

```
Module* makeLLVMModule() {  
    Module* mod = new Module("tut2");  
  
    // getOrInsertFunction() doesn't actually return a Function*  
    // it will return a cast of the existing function if the function  
    // already existed with a different prototype  
    Constant* c = mod->getOrInsertFunction("gcd",  
    /*ret type*/          IntegerType::get(32),  
    /*args*/              IntegerType::get(32),  
                          IntegerType::get(32),  
    /*varargs terminated with null*/ NULL);  
  
    Function* gcd = cast<Function>(c);  
}
```


makeLLVMModule()

```
// make looking at our output somewhat more
```

```
// pleasant
```

```
Function::arg_iterator args = gcd->arg_begin( );
```

```
Value* x = args++;
```

```
x->setName( "x" );
```

```
Value* y = args++;
```

```
y->setName( "y" );
```

makeLLVMModule()

```
// a block is basically a set of instructions that can be branched to  
// and is executed linearly until the block is terminated by one of a  
// small number of control flow instructions, such as br or ret  
// we're making use of LLVM's automatic name uniquing
```

```
BasicBlock* entry = BasicBlock::Create("entry", gcd);
```


```
BasicBlock* ret = BasicBlock::Create("return", gcd);
```

```
BasicBlock* cond_false = BasicBlock::Create("cond_false", gcd);
```

```
BasicBlock* cond_true = BasicBlock::Create("cond_true", gcd);
```

```
BasicBlock* cond_false_2 = BasicBlock::Create("cond_false", gcd);
```

CFG for 'gcd' function



makeLLVMModule()


IRBuilder<> builder(entry); // working on “entry” block

entry:

```
%tmp = icmp eq i32 %x, %y ; <i1> [#uses=1]  
br i1 %tmp, label %return, label %cond_false
```

```
Value* xEqualsY = builder.CreateICmpEQ(x, y, "tmp");  
builder.CreateCondBr(xEqualsY, ret, cond_false);
```

CFG for 'gcd' function



makeLLVMModule()

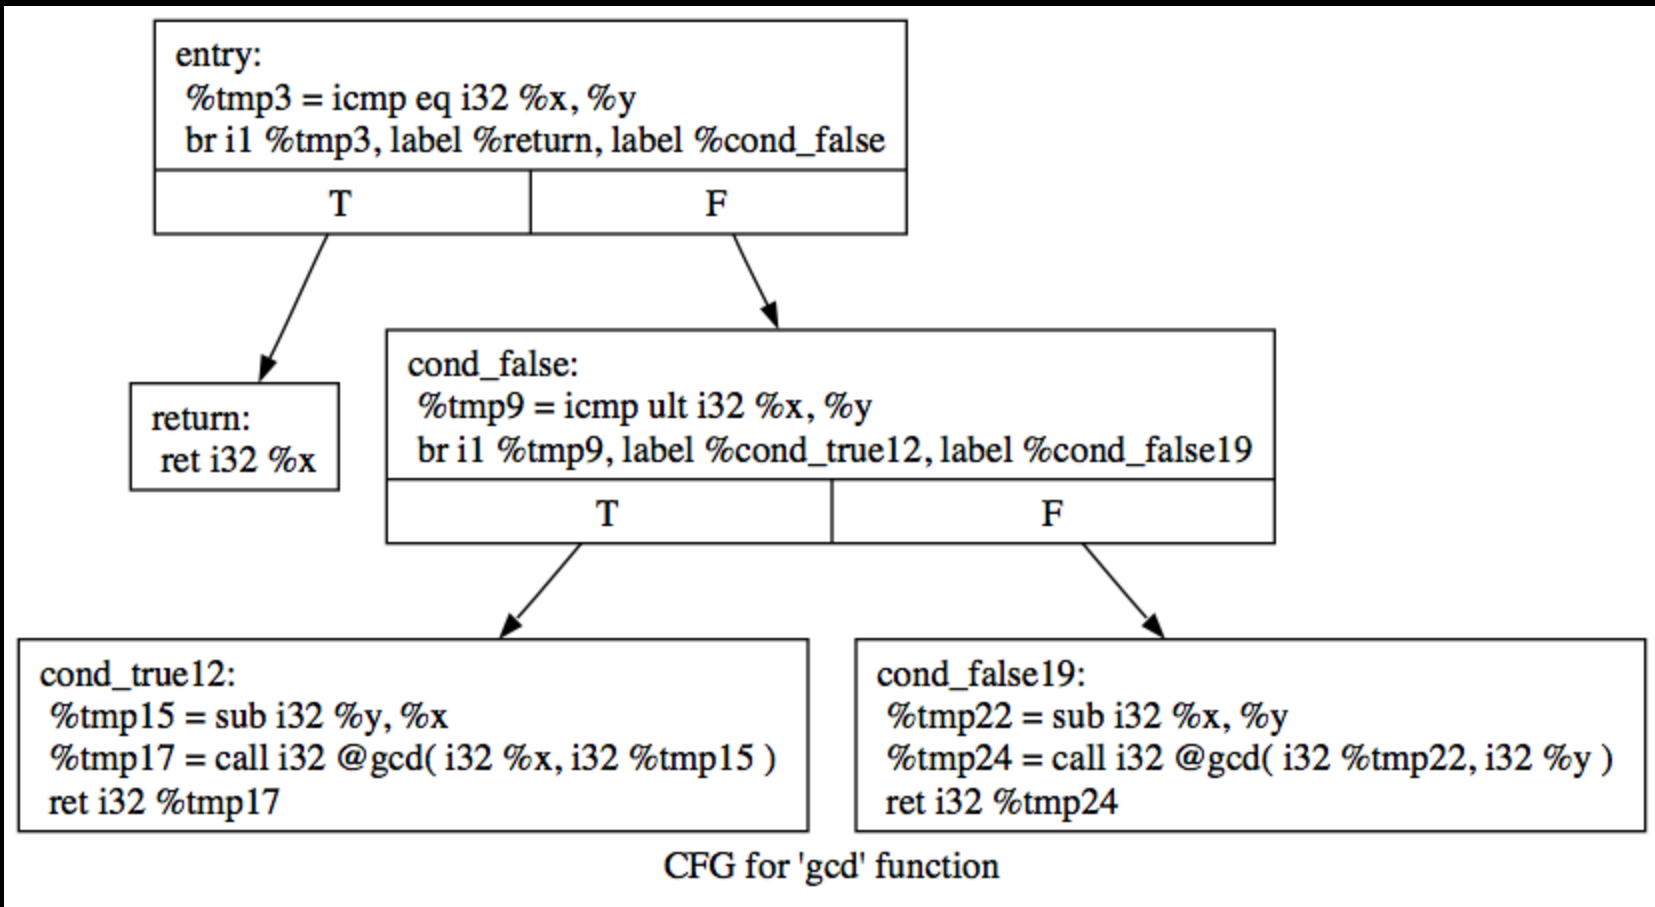
```
return:           ; preds = %entry  
    ret i32 %x
```

// use SetInsertPoint to retarget our current block

```
builder.SetInsertPoint(ret);
```

```
builder.CreateRet(x);
```

CFG for 'gcd' function



makeLLVMModule()

```
cond_false:                ; preds = %entry
    %tmp2 = icmp ult i32 %x, %y                ; <i1> [#uses=1]
    br i1 %tmp2, label %cond_true, label %cond_false1
```

// In LLVM, integer types do not carry sign

// Whether a signed or unsigned interpretation is desired is


// specified in the instruction

```
builder.SetInsertPoint(cond_false);
```

```
Value* xLessThanY = builder.CreateICmpULT(x, y, "tmp");
```

```
builder.CreateCondBr(xLessThanY, cond_true, cond_false_2);
```


CFG for 'gcd' function



makeLLVMModule()

```
cond_true:                ; preds = %cond_false
    %tmp3 = sub i32 %y, %x      ; <i32> [#uses=1]
    %tmp4 = call i32 @gcd(i32 %x, i32 %tmp3); <i32> [#uses=1]
    ret i32 %tmp4
```

// To create a call instruction, we have to create a vector (or any
// other container with InputIterators) to hold the arguments.

// We then pass in the beginning and ending iterators for this vector

```
builder.SetInsertPoint(cond_true);
```

```
Value* yMinusX = builder.CreateSub(y, x, "tmp");
```

```
std::vector<Value*> args1;
```


```
args1.push_back(x);
```

```
args1.push_back(yMinusX);
```

```
Value* recur_1 = builder.CreateCall(gcd, args1.begin(), args1.end(), "tmp");
```

```
builder.CreateRet(recur_1);
```

CFG for 'gcd' function



makeLLVMModule()

```
cond_false1:                ; preds = %cond_false
    %tmp5 = sub i32 %x, %y    ; <i32> [#uses=1]
    %tmp6 = call i32 @gcd(i32 %tmp5, i32 %y); <i32> [#uses=1]
    ret i32 %tmp6
```

```
Value* xMinusY = builder.CreateSub(x, y, "tmp");
```

```
std::vector<Value*> args2;
```

```
args2.push_back(xMinusY);
```

```
args2.push_back(y);
```

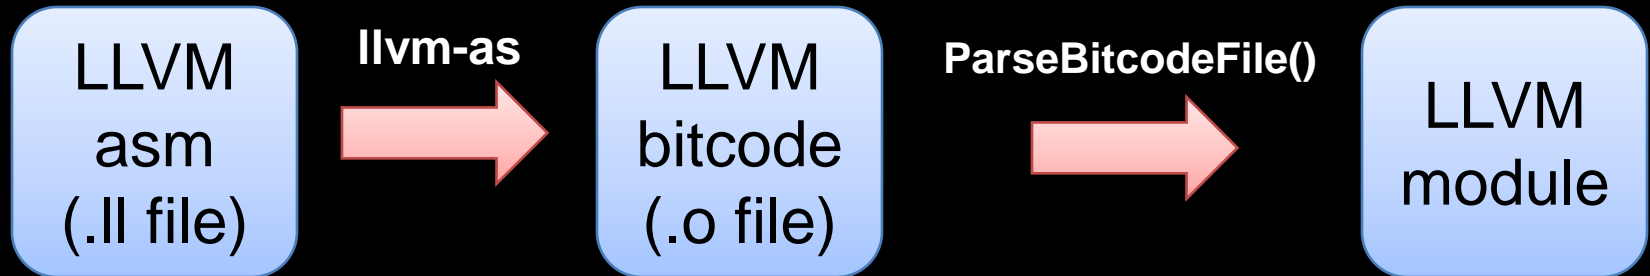
```
Value* recur_2 = builder.CreateCall(gcd, args2.begin(), args2.end(), "tmp");
```

```
builder.CreateRet(recur_2);
```

```
return mod;
```

```
}
```

Read LLVM Module



#include the appropriate LLVM header files

```
#include <llvm/Support/MemoryBuffer.h>
```

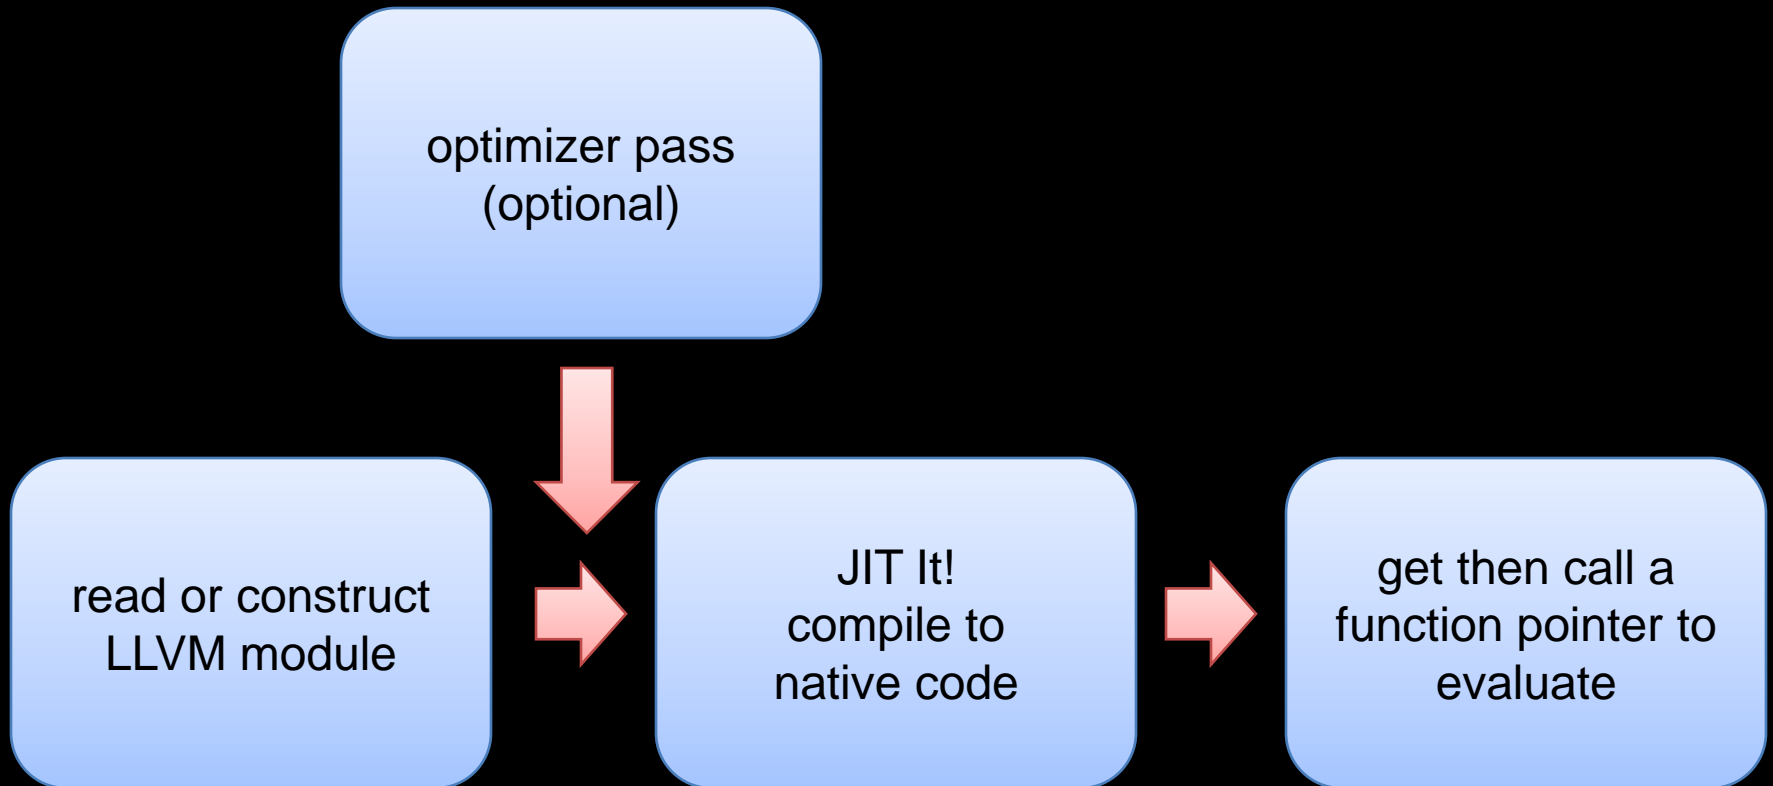
```
#include <llvm/Bitcode/ReaderWriter.h>
```

makeLLVMModule()

// create module from bitcode file

```
Module* makeLLVMModule( ) {  
    std::string error;  
    Module* Mod = 0;  
    // load in the bitcode file containing the functions  
    if( MemoryBuffer* buffer =  
        MemoryBuffer::getFile( "ops.o", &error ) ) {  
        // read the specified bitcode file, returning the module  
        Mod = ParseBitcodeFile( buffer, &error );  
        delete buffer;  
    }  
    return Mod;  
}
```

JIT & Execute



ExecutionEngine

- a key module to use JIT facility with LLVM
- can compile a function in the module into the native code
- provide a interface to call JIT-ted function

#include the appropriate LLVM header files

```
#include <llvm/ExecutionEngine/JIT.h>
```

```
#include <llvm/ModuleProvider.h>
```

```
#include <llvm/ExecutionEngine/GenericValue.h>
```

llvm/ExecutionEngine/JIT.h

- force the JIT to link in on certain operating systems (Windows)
- llvm/ExecutionEngine/ExecutionEngine.h do the real work
 - define the abstract interface that implements execution support for llvm

llvm/ModuleProvider.h

- an abstract interface for loading a module from some place
- allow incremental or random access loading of functions from the file

llvm/ExecutionEngine/GenericValue.h

- represent an LLVM value of arbitrary type

main()

```
// Now we create the JIT
```

```
// Create a ExecutinEngine for module Mod
```

```
ExistingModuleProvider* MP = new ExistingModuleProvider( Mod );
```

```
ExecutionEngine* EE = ExecutionEngine::create(MP, false);
```

```
// Get a (LLVM) function to be JIT-ed from module M
```

```
Function *GCD = cast<Function>( Mod->getOrInsertFunction("gcd",
    Type::Int32Ty,
    Type::Int32Ty,
    Type::Int32Ty,
    NULL) );
```

main()

// Call the function with arguments

```
std::vector<GenericValue> Args(2);
```

```
Args[0].IntVal = APInt(32, 10);
```

```
Args[1].IntVal = APInt(32, 5);
```

// Compile a function and execute it

// Provides function argument through GenericValue array

```
GenericValue gv = EE->runFunction(GCD, Args);
```

compile & run

```
$ c++ -g tut2.cpp `llvm-config --cxxflags --  
ldflags --libs engine` -o tut2
```

```
$ ./tut2
```


JIT Support

- JIT code generator emits machine code and auxiliary structures as binary output that can be written directly to memory
- currently Alpha, ARM, PowerPC, and X86 have JIT support

XXXCodeEmitter.cpp

- contain a machine function pass that transforms target-machine instructions (assembly) into relocatable machine code (binary)

```
// relocate function in XXXJITInfo.cpp will fill in the actual address  
// of SYMBOL  
jump    SYMBOL  
        ... (omit)  
SYMBOL
```

XXXJITInfo.cpp

- implement the JIT interfaces for target-specific code-generation activities
- minimum implementation
 - getLazyResolverFunction
 - emitFunctionStub
 - relocate
 - callback

getLazyResolverFunction

- initialize the JIT, gives the target **a function that is used for compilation**
- trivial implementation - make the parameter as the global **JITCompilerFunction** and returns the callback function that will be used a function wrapper

```
// ARMJITInfo.cpp
```

```
TargetJITInfo::LazyResolverFn
```

```
ARMJITInfo::getLazyResolverFunction(JITCompilerFn F) {
```

```
    JITCompilerFunction = F;
```

```
    return ARMCompilationCallback;
```

```
}
```

http://llvm.org/doxygen/ARMJITInfo_8cpp-source.html

callback

```
// ARMJITInfo.cpp
void ARMCompilationCallback(void); // the callback as a wrapper
asm(
    //store register before calling JIT
    "stmdb sp!, {r0, r1, r2, r3, lr}\n"

    "mov r0, lr\n"
    "sub sp, sp, #4\n"
    // call the C portion of the callback
    "bl " ASMPREFIX "ARMCompilationCallbackC\n"
    "add sp, sp, #4\n"
    // restore register after calling JIT
    "ldr r0, [sp,#20]\n"
    "ldr r1, [sp,#16]\n"
    "str r1, [sp,#20]\n"
    "str r0, [sp,#16]\n"
    "ldmia sp!, {r0, r1, r2, r3, lr, pc}\n"
```

emitFunctionStub

```
// ARMJITInfo.cpp
// returns a native function with a specified address for a callback function
void *ARMJITInfo::emitFunctionStub(const Function* F, void *Fn,
                                   MachineCodeEmitter &MCE) {
    // If this is just a call to an external function, emit a branch instead of a
    // call. The code is the same except for one bit of the last instruction.
    if (Fn != (void*)(intptr_t)ARMCompilationCallback) {
        // branch to the corresponding function addr
        // the stub is 8-byte size and 4-aligned
    } else {
        // branch and link to the compilation callback
        // the stub is 16-byte size and 4-aligned
    }
    ... (omit)
}

void ARMCompilationCallback(void);
```

relocate

```
// ARMJITInfo.cpp

// changes the addresses of referenced globals, based on relocation types
/// relocate - Before the JIT can run a block of code that has been emitted,
/// it must rewrite the code to contain the actual addresses of any
/// referenced global symbols.

void ARMJITInfo::relocate(void *Function, MachineRelocation *MR,
                          unsigned NumRelocs, unsigned char* GOTBase) {
    for (unsigned i = 0; i != NumRelocs; ++i, ++MR) {
        ...(omits)
    }
    switch ((ARM::RelocationType)MR->getRelocationType()) {
        case ARM::reloc_arm_relative: {
        }
        case ARM::reloc_arm_branch: {
        }
    }
}
```

XXXTargetMachine

- add **getJITInfo** method that return a TargetJITInfo object

```
// AlphaTargetMachine.h  
  
virtual AlphaJITInfo* getJITInfo( ) {  
    return &JITInfo;  
}
```


References

- LLVM Tutorial
- LLVM Language Reference Manual
- LLVM - 2.0 and beyond!
- Writing an LLVM Pass
- Writing an LLVM Compiler Backend
- Adding LLVM JIT – Syoyo Fujita
- http://llvm.org/devmtg/2008-08-23/llvm_first_steps.pdf
- <http://blogger.godfat.org/2008/10/llvm-rubinius.html>
- `$(LLVM_SRC_ROOT)/examples/HowToUseJIT`
- `$(LLVM_SRC_ROOT)/lib/Target`